Rahul Muralidhara Koushik

San Francisco Bay Area
Email: rkoushik@asu.edu
Mobile: +1-480-791-6597
GitHub: github.com/Fluff18 | Portfolio: fluff18.github.io
LinkedIn: linkedin.com/in/rahulkoushik

AI Engineer | LLM Agents, Evals, Inference. Builds generative AI (GenAI) and large language model (LLM) agents end to end: the user problem, the agent loop, the memory layer and the evals that decide whether an agent is ready to ship. Three HPE hackathon wins, two selected for Fellow Review and BU-wide rollout; HPE Early Career Technologist. Interested in AI engineering, AI product and physical AI.

Experience

Hewlett Packard Enterprise, Storage AI Team — San Jose, US
AI Engineer — *Feb 2026 – Present*

- Architected the production **multi-agent LLM harness** with **guardrails** that every team agent runs on; built and maintain its **MCP** server and tools, used across HPE Storage software engineering (**6,300+ tool calls a week**).
- Built the **memory layer** for persistent agents: episodic memory, **embeddings and vector search**, and a **context graph** over **13,005** mined bugs and **5,600** grounded analyses.
- Cut agent latency **45%** with **cache-coherent** context updates, plus **prompt caching**, **KV-cache reuse** and **prefill/decode-aware** serving of **Qwen** (27B, 35B) on **vLLM**.
- Took the **agent evaluation platform** from idea to **three production deployments in four weeks**; rubric-based evals, **LLM-as-a-judge** scoring and a **tracing dashboard** for every agent run form the **release gate**.
- Shipped **agentic alert-triage and root-cause-analysis (RCA)** systems to production, with **self-improving loops**: a reflector agent reviews each run's path and SMEs confirm the fixes.

Hewlett Packard Enterprise, Storage — San Jose, US
Cloud Software Engineer — *Jun 2023 – Feb 2026*

- Built the **observability platform** for internal teams supporting **10K customers**: cross-stack root-cause analysis cut resolution time from **30 to 3 minutes** and raised first-time fix success from **55% to 85%**.
- Shipped **6 enterprise data-migration features** in Go across four generations of storage arrays, over Fibre Channel and iSCSI.
- Built a **documentation agent** that updates docs on every pull request, saving **5 engineer-hours a week** per team.

The Luminosity Lab, Arizona State University — Tempe, US
Research Aide — *Feb 2022 – May 2023*

- Led the iOS team on the TouchPoint apps; a telehealth flow saved doctors **15 minutes** per visit.

Aruba, Hewlett Packard Enterprise — Tempe, US
Software Engineering Intern — *May 2022 – Aug 2022*

- Built a test-automation tool that configures network topologies, saving an estimated **150k engineering hours a year**.

Webex by Cisco — Bengaluru, India
Technical Undergraduate Intern — *Jan 2021 – Jul 2021*

- Built call and meeting scheduling APIs for the Webex iOS SDK v3 and the Kitchen Sink sample app.

Skills

• **Generative AI (GenAI)**: LLM agents, agentic workflows, MCP and tool calling, evals and LLM-as-a-judge, Retrieval-Augmented Generation (RAG), embeddings and vector search (pgvector), guardrails, multi-agent orchestration, prompt and context engineering, memory, tracing, LangGraph, vLLM, Qwen, llama.cpp

• **Languages**: Python, Go, TypeScript, JavaScript, Swift, Kotlin, C++, SQL

• **Infrastructure**: Distributed systems, Kubernetes, Helm, Docker, AWS, Kafka, Spark, Postgres, Prometheus, Grafana, gRPC, FastAPI, React, Next.js

Selected Projects

• **Zuo, personal AI harness (2026)**: Agent teams, job queue and budgets run from my phone across three machines.

• **Catan Arena, self-learning agents (2026)**: Agents graduate only when a statistical pass bar says the learning is real.

• **Zzz, iOS and Android (2026)**: Gamified goal tracker, live on the App Store. SwiftUI, Kotlin, Supabase. App Store

Education and Publications

M.S. Computer Science, Arizona State University, Tempe — Aug 2021 – May 2023

B.Tech Computer Science and Engineering, PES University, Bengaluru — Aug 2017 – Jul 2021

Publications: ViStA: Video Streaming and Analytics Benchmark, IEEE ISPASS 2021; Generation of True Random Numbers using Entropy Sources in Portable Computers, IEEE CONECCT 2021; PipeGAN, a pipelined parallel GAN, Springer LNNS vol. 171, 2021.